

# Intro to NoSQL

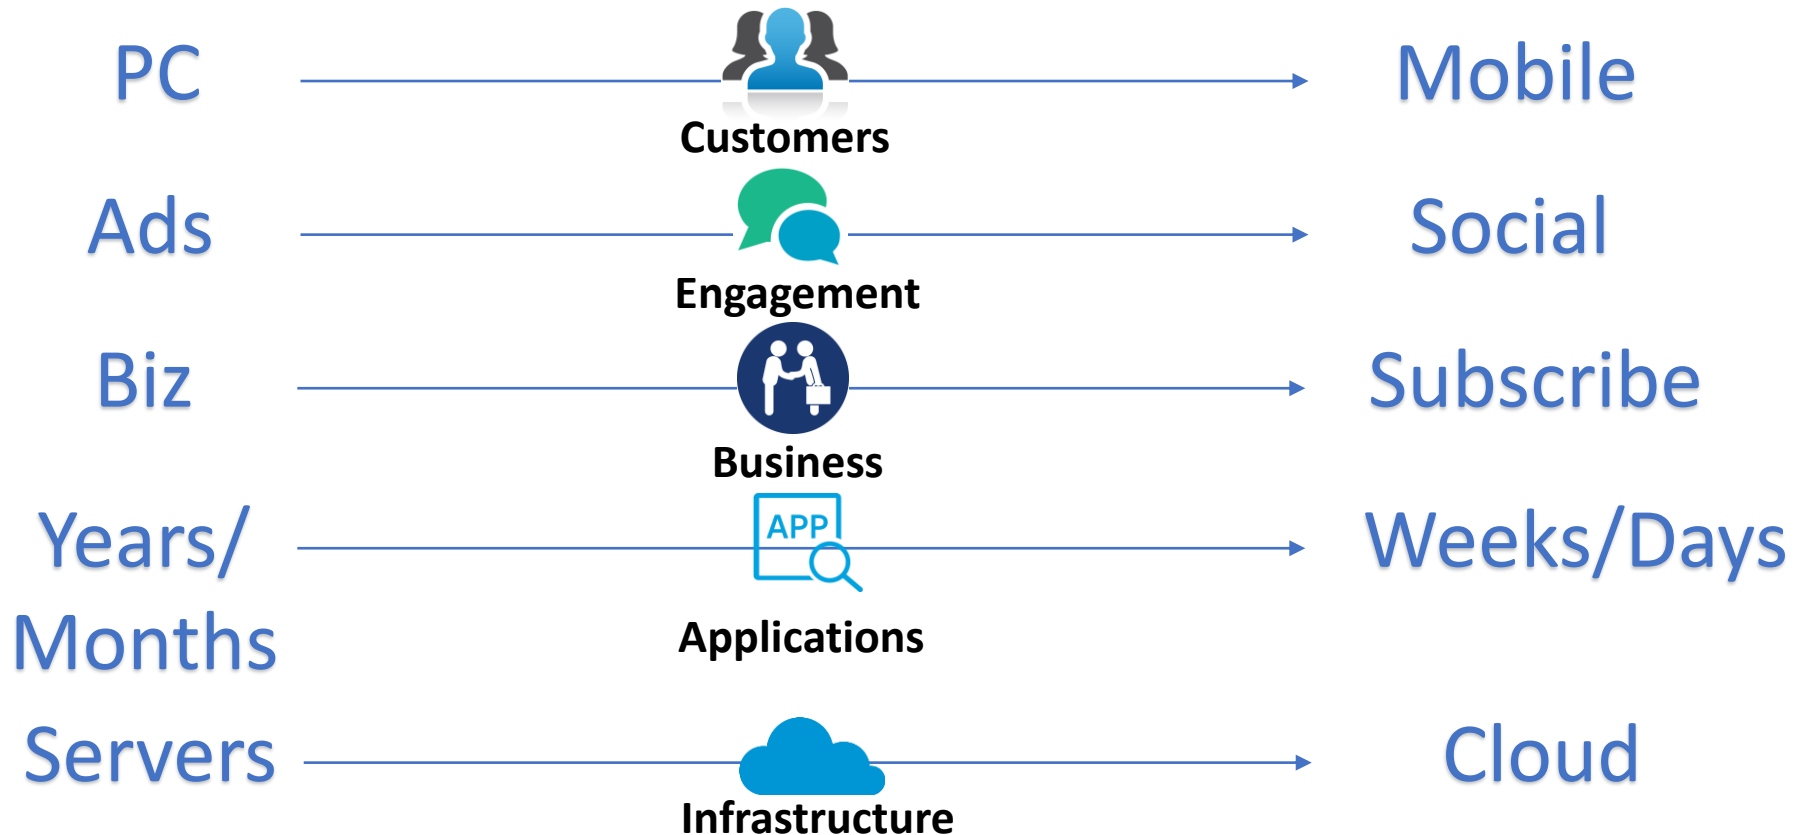
Poo Kuan Hoong

# Introduction

- A NoSQL (originally referring to "non SQL" or "non relational") database provides a mechanism for storage and retrieval of data that is modeled in means other than the tabular relations used in relational databases.

# Digital platforms have changed

- Digital platforms have changed over the past 5 years




# The New Enterprise Stack

	Traditional	Modernized
Apps	On-premised, monoliths	SaaS, Microservices
Database	Relational (Oracle)	Non-relational (MongoDB)
EDW	Teradata, Oracle, etc	Hadoop
Compute	Scale-up server	Containers/Commodity Server/ Cloud
Storage	SAN	Local Storage & Data Lakes
Network	Routers & Switches	Software defined Networks

# Why do we need NoSQL?

- We are storing more data now than we ever have before
- Connections between our data are growing all the time
- We don't make things knowing the structure from day 1
- Server architecture is now at a stage where we can take advantage of it




# The challenges of using RDBMS

## Legacy

### Relational Model

Rigid Schemas  
Resistant to  
change



Data changes constantly,  
which fits poorly with a  
relational model

### Scale-up

Throughput & Cost  
make Scale-Up  
Impractical



Scale-Up clusters were never  
meant to handle today's  
volumes

## Today

### Flexible Model



Flexible Multi-Structured  
Schema that is designed to  
adapt to changes

### Scale-out



Scale-out to the end of the world  
and distribute data where it  
needs to be

# NoSQL Use Cases

- Large data volumes
  - Massively distributed architecture
  - Required to store the data
  - Google, Amazon, Facebook, 100k servers
- Extreme query workload
  - Impossible to efficiently do joins at that
  - Scale with an RDBMS
- Schema evolution
  - Schema flexibility is not trivial at a large
  - Scale but it can be with NoSQL

# NoSQL: Pros and Cons

- **Pros**

- Massive scalability
- High availability
- Lower cost
- Schema flexibility
- Sparse and semi structured data

- **Cons**

- Limited query capabilities
- Not standardised (portability may be an issue)
- Still a developing technology



**BIGTABLE**

**KEY VALUE**

**FOUR**

EMERGING TRENDS IN  
NOSQL DATABASES

**GRAPHDB**

**DOCUMENT**

# Big Table

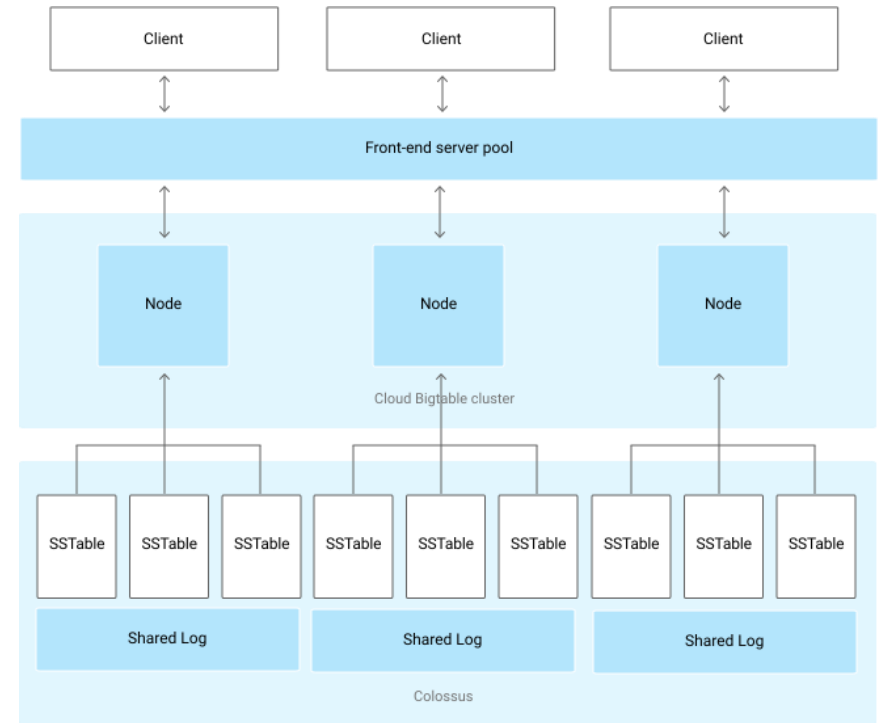
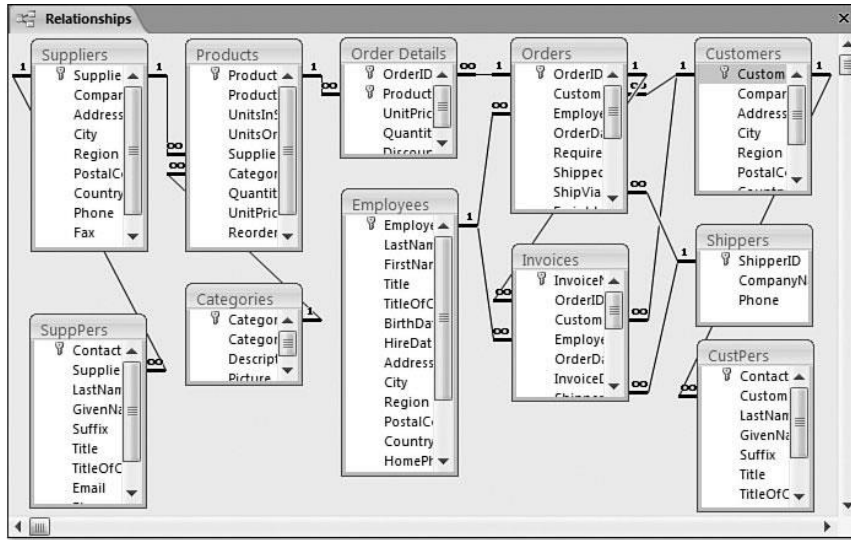
- Behaves like a standard relational database but with a slight change
- Designed to work with a lot of data...a really big crap ton
- Created by google and now used by lots of others
- <http://research.google.com/archive/bigtable.html>
- <http://research.google.com/archive/spanner.html>

Google  
BigTable



# Big Table

- “A Bigtable is a sparse, distributed, persistent multidimensional sorted map. The map is indexed by a row key, column key, and a timestamp; each value in the map is an uninterpreted array of bytes.”



# Relational Database vs Big Table

# Key value database

- Again, designed to work with a lot of data
- Each bit of data is stored in a single collection
- Each collection can have different types of data



# Document Oriented Database

- A document-oriented database, or document store, is a computer program designed for storing, retrieving and managing document-oriented information, also known as semi-structured data. Document-oriented databases are one of the main categories of NoSQL databases
- Very similar to a key value database, where main difference is that you can actually see the values
- CRUD operations

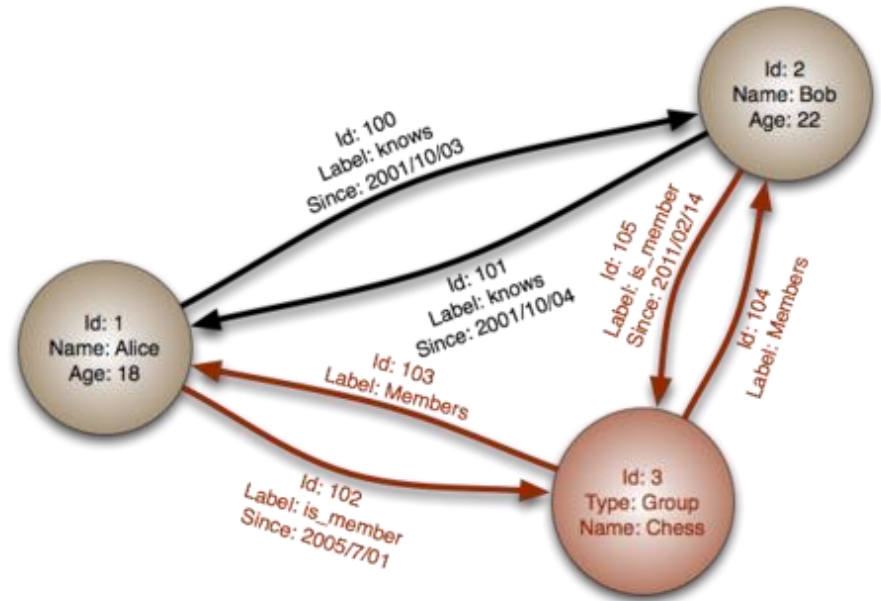


mongoDB



# Graph Database

- Focus here is on modelling the structure of the data
- Inspired by graph theory
- Scales really well to the structure of the data



# Graph Database

- [Neo4j](#) is a graph database management system developed by Neo4j, Inc. Described by its developers as an ACID-compliant transactional database with native graph storage and processing
- Neo4j is the most popular graph database according to DB-Engines ranking





# NoSQL: Tips

- High availability and disaster recovery are a must
- Understand the pros and cons of each design model
- Don't pick something just because it is new
- Don't pick something based JUST on performance

# SQL: The good

- High performance for transactions. Think [ACID](#)
- Highly structured, very portable
- Small amounts of data
- SMALL IS LESS THAN 500GB
- Supports many tables with different types of data
- Can fetch ordered data
- Compatible with lots of tools



# SQL: ACID

- Atomicity
- Consistency
- Isolation
- Durability

# SQL: The bad

- Complex queries take a long time
- The relational model takes a long time to learn
- Not really scalable
- Not suited for rapid development



# NoSQL : The good

- Fits well for volatile data
- High read and write throughput
- Scales really well
- Rapid development is possible
- In general it's faster than SQL



# NoSQL: The bad

- Key/Value pairs need to be packed/unpacked all the time
- Still working on getting security for these working as well as SQL
- Lack of relations from one key to another



# Summary

## SQL

works great, can't scale for large data

## NoSQL

works great, doesn't fit all situations

so use both, but think about when you want to use them!